



FIGURE 3.27 Floating-point format for IEEE 754 single precision (fp32), IEEE 754 half-precision (fp16), and Brain Float 16. Google’s TPUv3 hardware uses Brain Float 16 (see Section 6.11).