

```

1  #include <x86intrin.h>
2  #define UNROLL (4)
3  #define BLOCKSIZE 32
4  void do_block (int n, int si, int sj, int sk,
5                double *A, double *B, double *C)
6  {
7      for ( int i = si; i < si+BLOCKSIZE; i+=UNROLL*8 )
8          for ( int j = sj; j < sj+BLOCKSIZE; j++ ) {
9              __m512d c[UNROLL];
10             for (int r=0;r<UNROLL;r++)
11                 c[r] = _mm512_load_pd(C+i+r*8+j*n); //[ UNROLL];
12
13             for( int k = sk; k < sk+BLOCKSIZE; k++ )
14             {
15                 __m512d bb = _mm512_broadcastsd_pd(_mm_load_sd(B+j*n+k));
16                 for (int r=0;r<UNROLL;r++)
17                     c[r] = _mm512_fmadd_pd(_mm512_load_pd(A+n*k+r*8+i), bb, c[r]);
18             }
19
20             for (int r=0;r<UNROLL;r++)
21                 _mm512_store_pd(C+i+r*8+j*n, c[r]);
22         }
23     }
24
25 void dgemm (int n, double* A, double* B, double* C)
26 {
27     for ( int sj = 0; sj < n; sj += BLOCKSIZE )
28         for ( int si = 0; si < n; si += BLOCKSIZE )
29             for ( int sk = 0; sk < n; sk += BLOCKSIZE )
30                 do_block(n, si, sj, sk, A, B, C);
31 }

```

FIGURE 5.48 Optimized C version of DGEMM from Figure 4.80 using cache blocking. These changes are the same ones found in Figure 5.21. The assembly language produced by the compiler for the `do_block` function is nearly identical to Figure 4.81. Once again, there is no overhead to call the `do_block` because the compiler inlines the function call.